

CS 580: Algorithm Design and Analysis

Upper and Lower Bounds for Sorting Algorithms

Merge Sort: $O(n * \log(n))$

Q: Can we do better?

Answer: It depends on the model of computation.
Comparisons counted as the expensive operation

Which is which?



Credit to Mary Wooters for lower bound slides


Comparison-based sorting



Want to sort these items.

There's some ordering on them, but we don't know what it is.

Is  bigger than  ?



Algorithm

YES

The algorithm's job is to
output a correctly sorted
list of all the objects.



There is a **genie** who knows what
the right order is.

The genie can answer YES/NO
questions of the form:
is [this] bigger than [that]?

Merge Sort and many other sorting algorithms work like this.



Is 7 bigger than 5 ?

YES

Is 6 bigger than 5 ?


YES


Is 3 bigger than 5 ?

NO



etc.






Lower bound of $\Omega(n \log(n))$.

- Theorem:
 - Any deterministic comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps.

This covers all the
sorting algorithms
we know!!!

- How might we prove this?
 1. Consider all comparison-based algorithms, one-by-one, and analyze them.



Lower bound of $\Omega(n \log(n))$.

- Theorem:

- Any deterministic comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps.

This covers all the
sorting algorithms
we know!!!

- How might we prove this?

1. Consider all comparison-based algorithms, one-by-one, and analyze them.

2. Don't do that.

Instead, argue that all comparison-based sorting algorithms give rise to a **decision tree**.

Then analyze decision trees.

Decision trees




Sort these three things.



YES

NO

etc...



YES

NO



YES

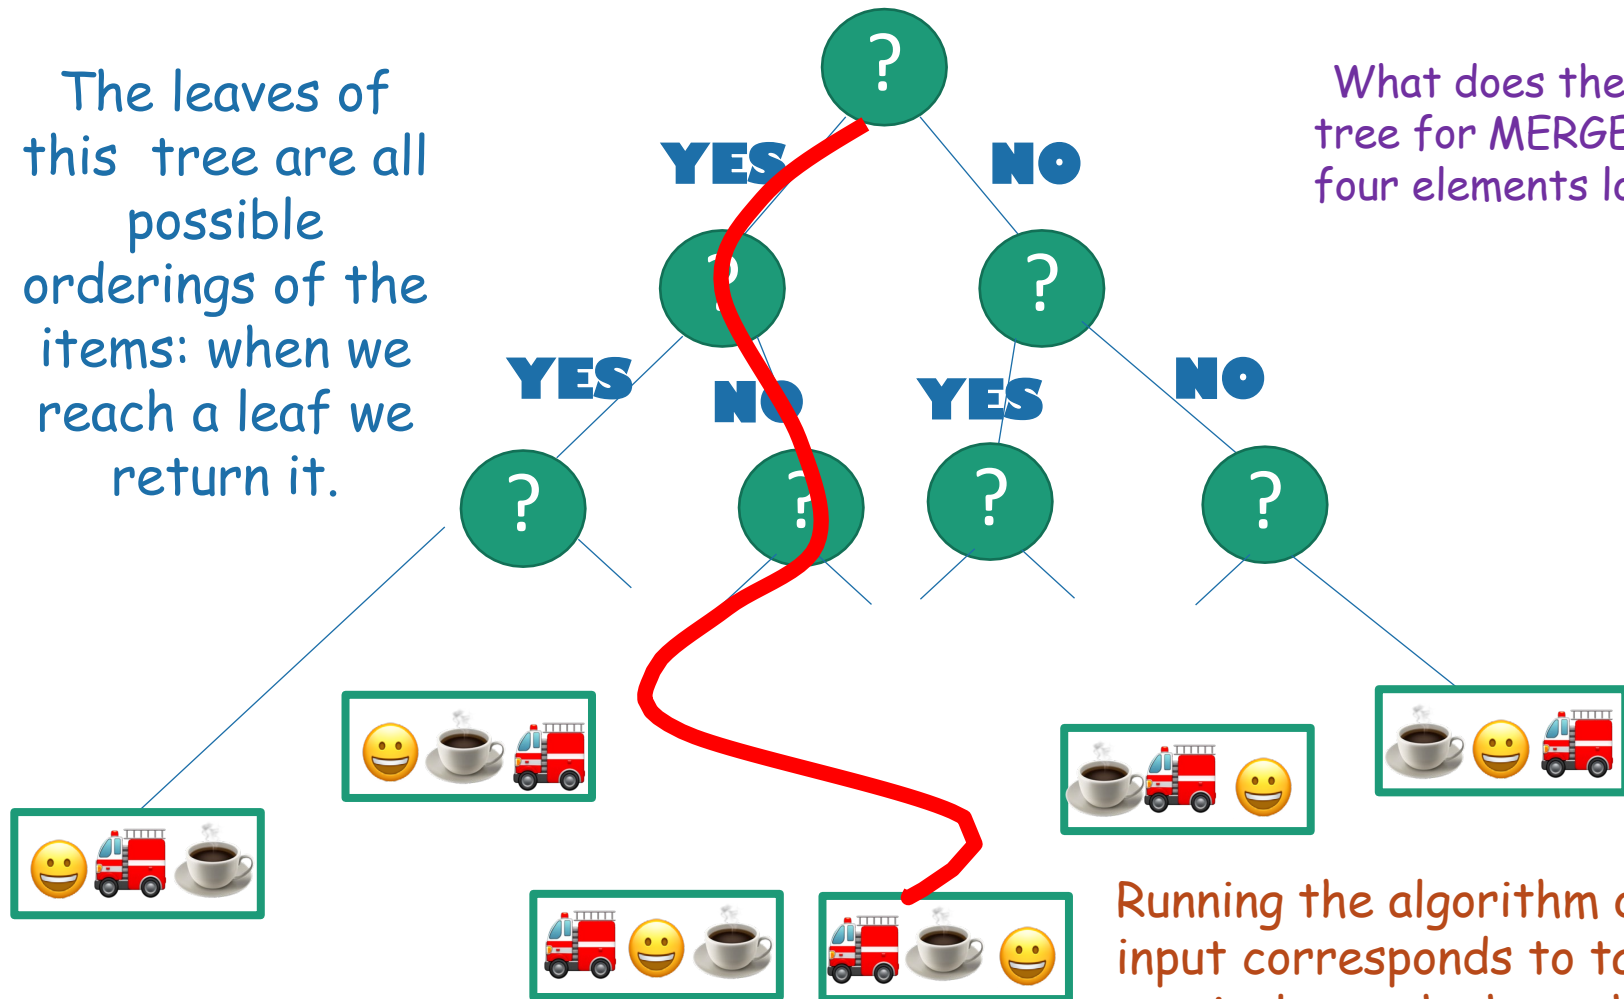
NO



All comparison-based algorithms have an associated decision tree.

The leaves of this tree are all possible orderings of the items: when we reach a leaf we return it.

What does the decision tree for MERGESORTING four elements look like?

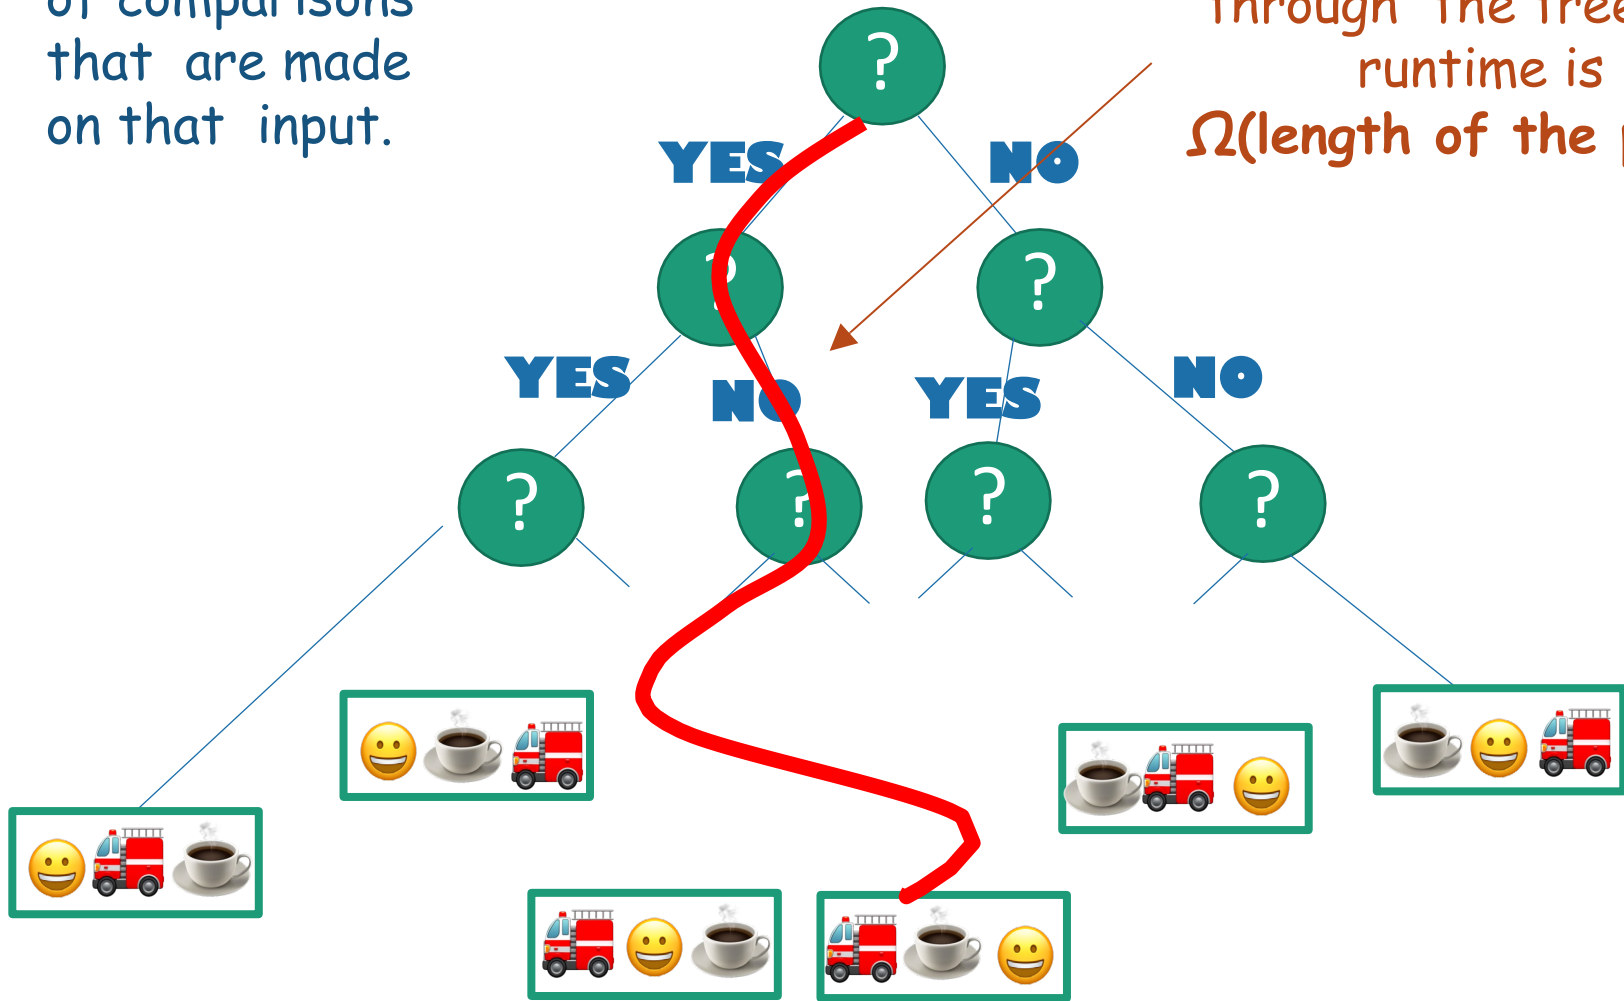


Running the algorithm on a given input corresponds to taking a particular path through the tree.

What's the runtime on a particular input?

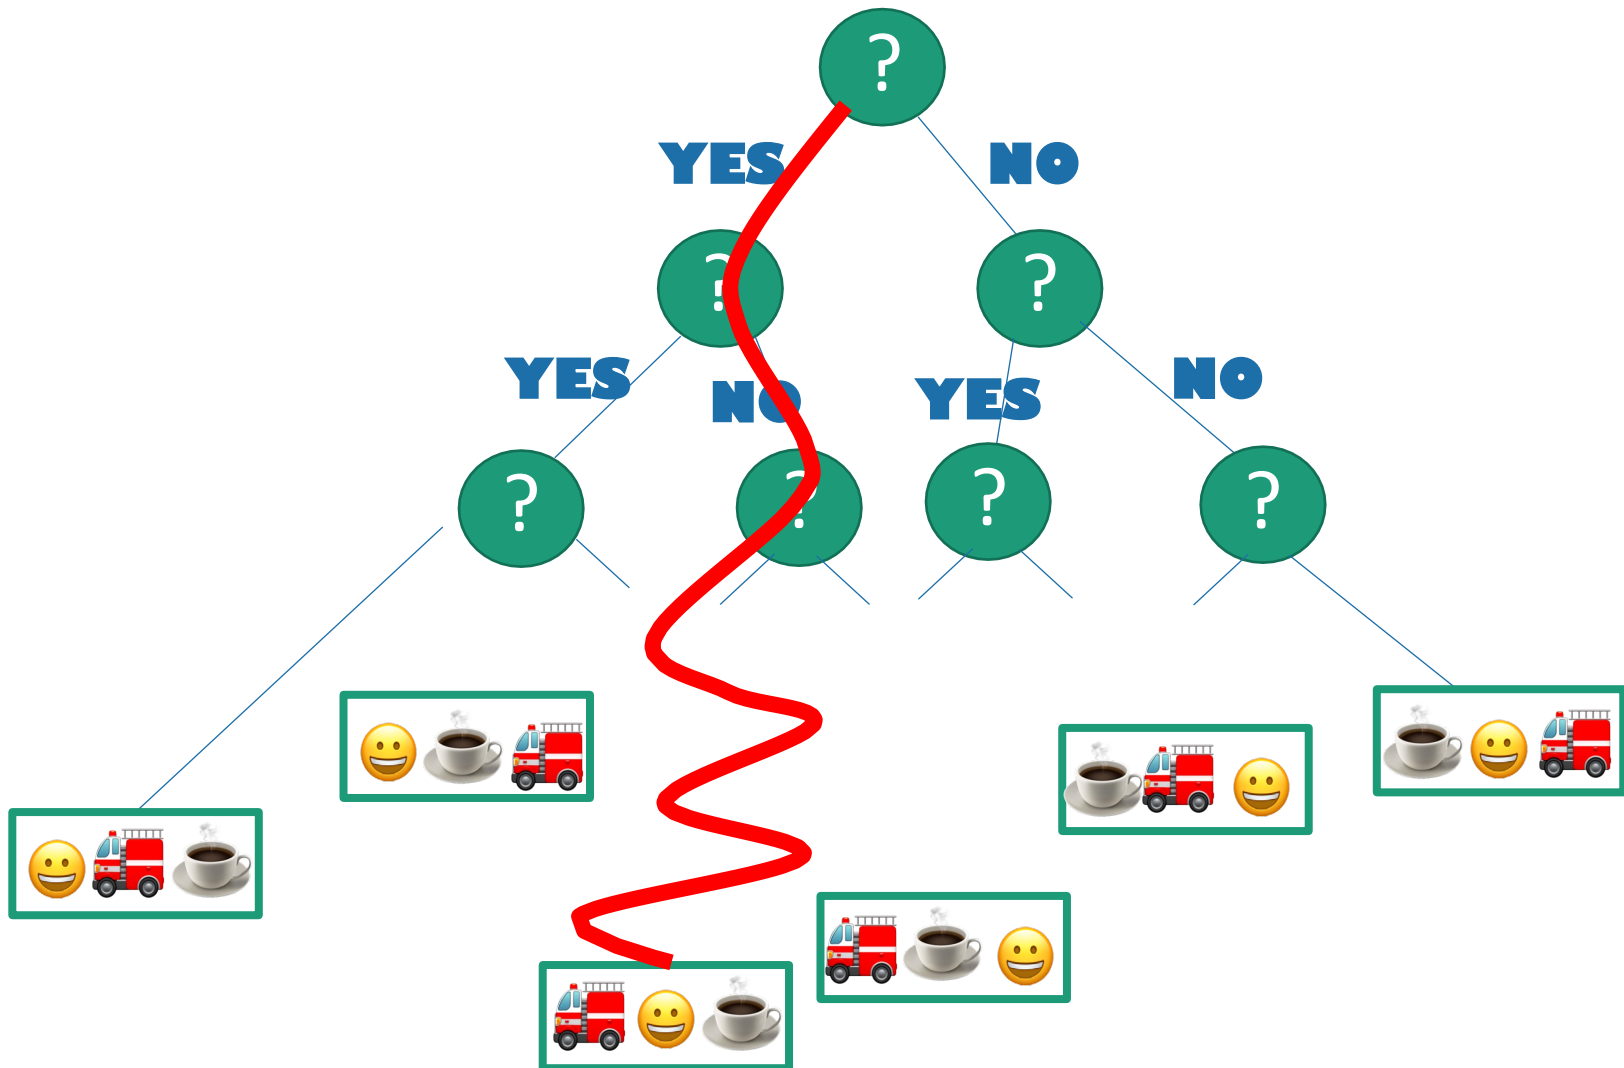
At least the number
of comparisons
that are made
on that input.

If we take this path through the tree, the runtime is $\Omega(\text{length of the path})$.



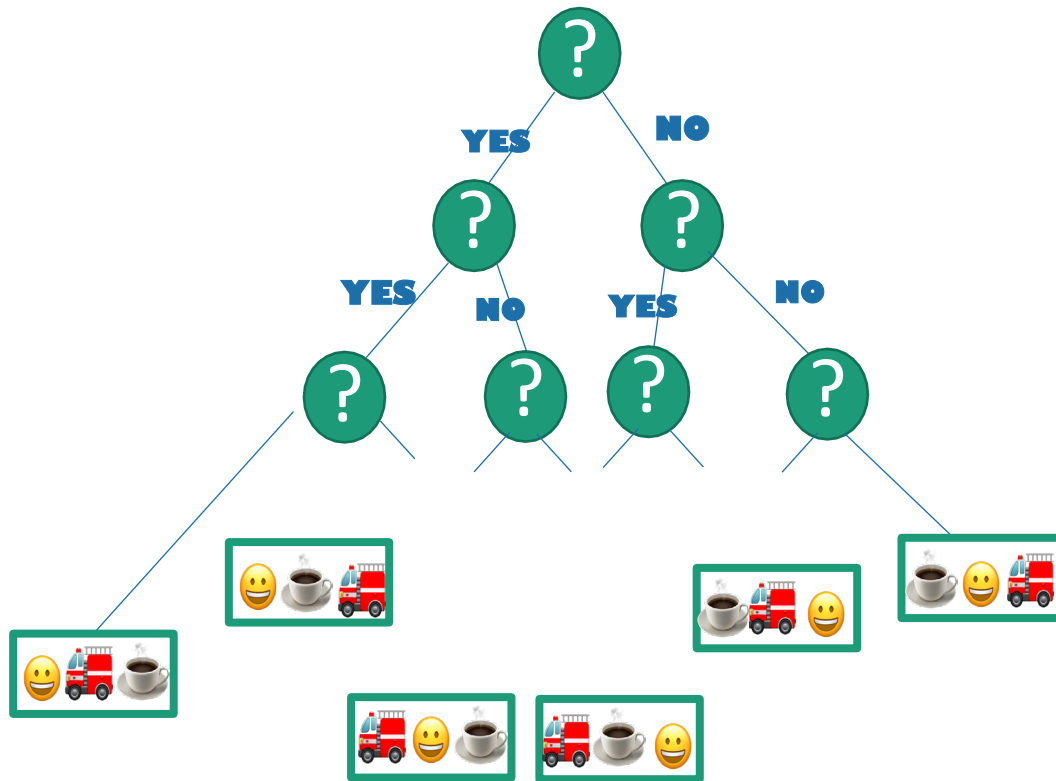
What's the worst-case runtime?

At least $\Omega(\text{length of the longest path})$.



How long is the longest path?

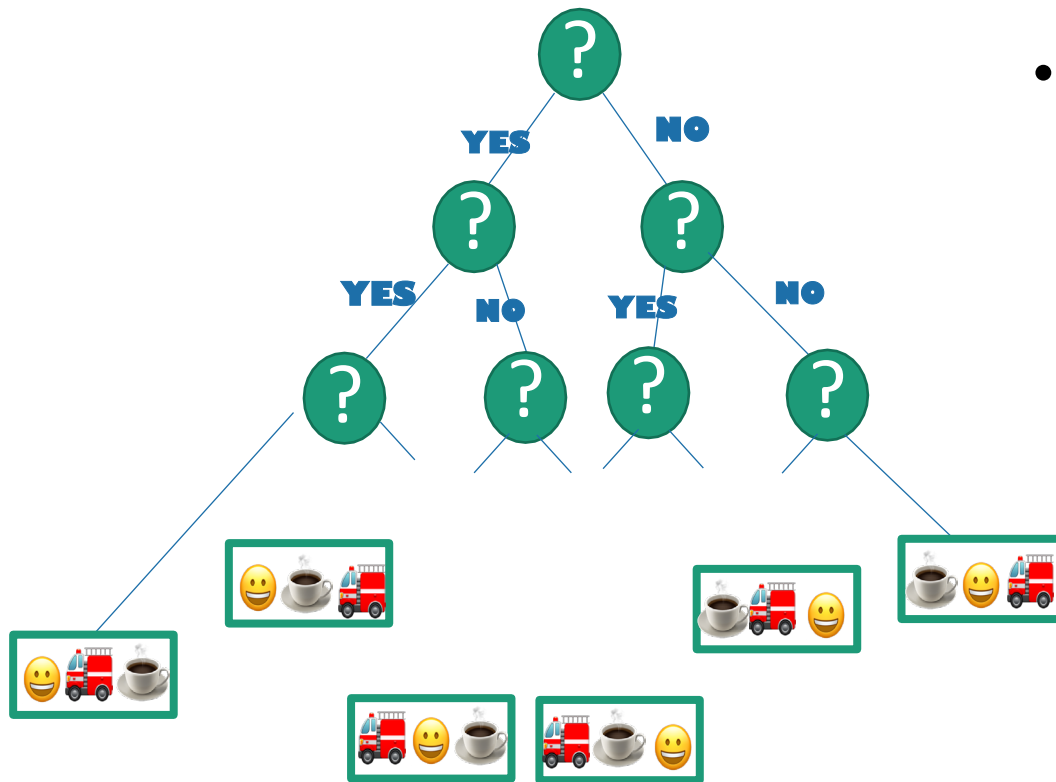
We want a statement: in all such trees, the longest path is at least _____



How long is the longest path?

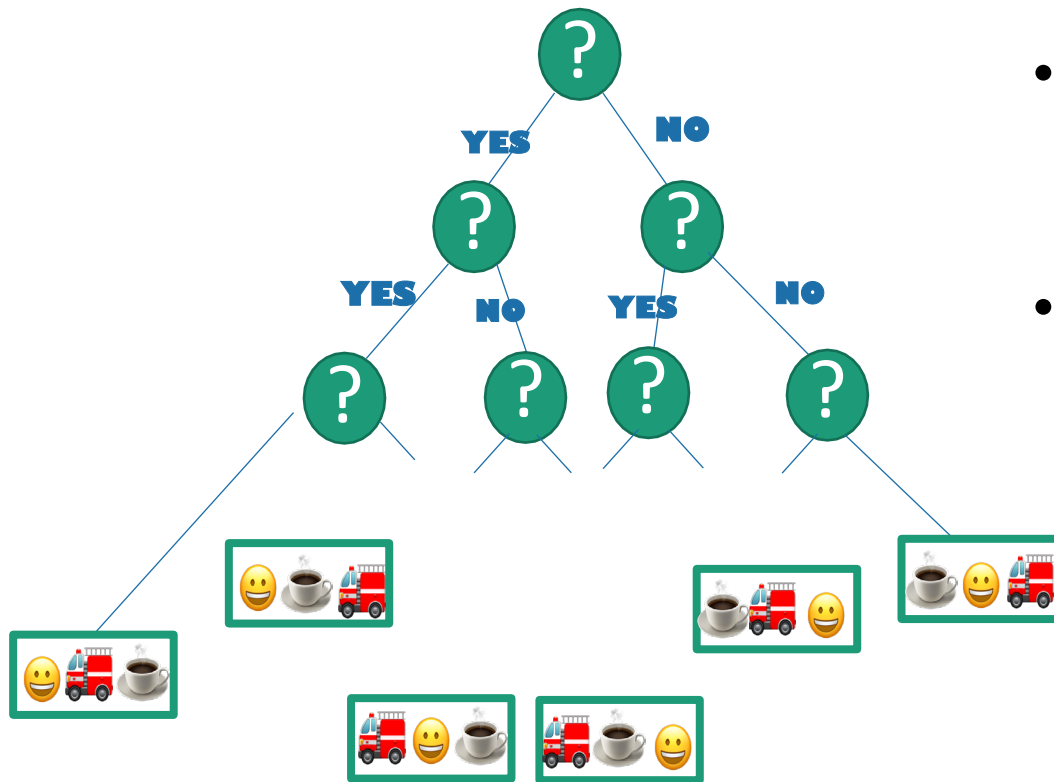
We want a statement: in all such trees, the longest path is at least _____

- This is a binary tree with at least _____ leaves.



How long is the longest path?

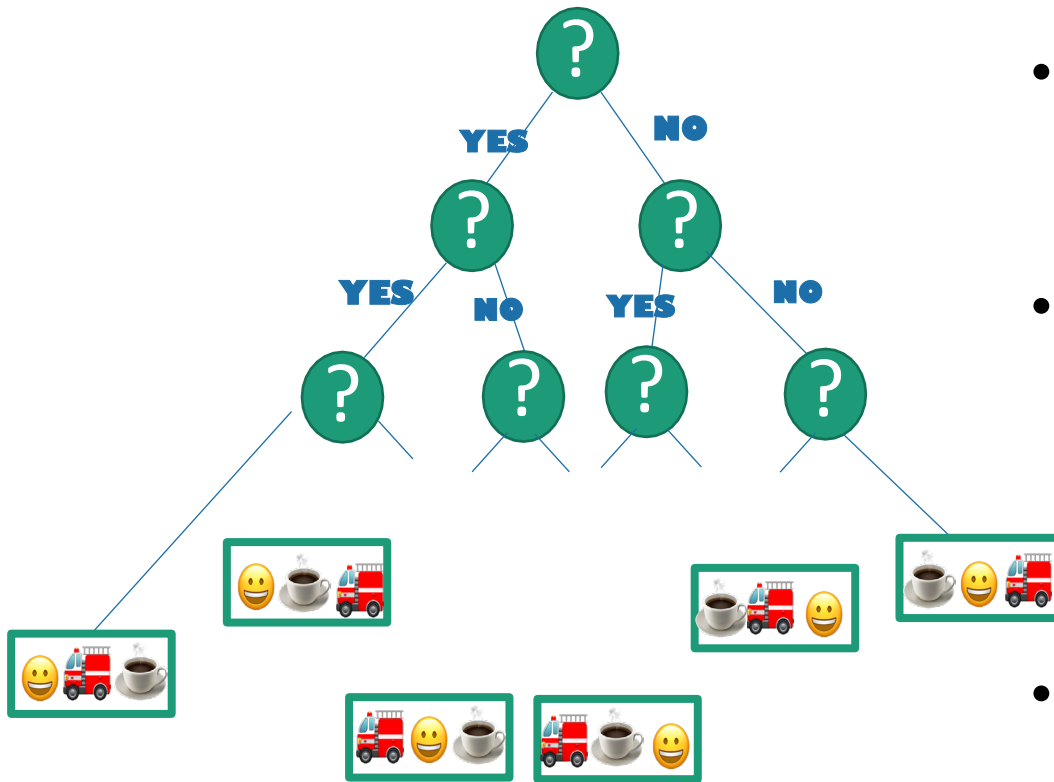
We want a statement: in all such trees, the longest path is at least _____



- This is a binary tree with at least $n!$ leaves.
- The shallowest tree with $n!$ leaves is the completely balanced one, which has depth _____

How long is the longest path?

We want a statement: in all such trees, the longest path is at least _____




- This is a binary tree with at least $n!$ leaves.
- The shallowest tree with $n!$ leaves is the completely balanced one, which has depth $\log(n!)$.
- So in all such trees, the longest path is at least $\log(n!)$.

- $n!$ is about $(n/e)^n$ (Stirling's formula).
- $\log(n!)$ is about $n \log(n/e) = \Omega(n \log(n))$.

Conclusion: the longest path has length at least $\Omega(n \log(n))$.

Lower bound of $\Omega(n \log(n))$.



- Theorem:
 - Any deterministic comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps.
- Proof:
 - Any deterministic comparison-based algorithm can be represented as a decision tree with $n!$ leaves.
 - The worst-case running time is at least the depth of the decision tree.
 - All decision trees with $n!$ leaves have depth $\Omega(n \log(n))$.
 - So any comparison-based sorting algorithm must have worst-case running time at least $\Omega(n \log(n))$.